

## > Set Up

↳ 1 cell hidden

## ✓ Data Preparation

```
# Locate the first row of 2019
```

```
# Seperate the dataset into smaller part to display
```

```
chunk_size = 700000
```

```
chunks = pd.read_csv('Combined Listing Data [Detailed].csv', chunksize=chunk_size)
```

```
df = next(chunks)
```

```
# # Create mask for rows in year 2019
```

```
df["year"] = pd.to_datetime(df['calendar_last_scraped'], errors="coerce").dt.year
```

```
mask_2019 = df['year'] == 2019
```

```
# # Grab the first row that meets the condition
```

```
first_2019_index = df[mask_2019].index[0]
```

```
print(first_2019_index)
```

```
<ipython-input-2-c0da483c02f7>:6: DtypeWarning: Columns (0,1,13,14,16,17,25,38,41,46,51,53,54,65,66,68,71,72,74,85,86,93,
df = next(chunks)
628693
```

```
# Read data and clean data
```

```
CSV_FILE = 'Combined Listing Data [Detailed].csv'
```

```
start_row = 628694 # inclusive
```

```
chunk_size = 500000 # number of rows per chunk
```

```
wanted_columns = [
```

```
    'calendar_last_scraped',
```

```
    'review_scores_accuracy',
```

```
    'review_scores_cleanliness',
```

```
    'review_scores_checkin',
```

```
    'review_scores_communication',
```

```
    'review_scores_location',
```

```
    'review_scores_value',
```

```
    'host_response_rate',
```

```
    'host_is_superhost',
```

```
    'instant_bookable',
```

```
    'description',
```

```
    'cleaning_fee'
```

```
]
```

```
# We will keep reading until we can no longer get any rows
```

```
all_chunks = [] # list to store each chunk's DataFrame
```

```
current_skip = start_row
```

```
while True:
```

```
    # skiprows=range(1, current_skip) means skip lines 1..(current_skip-1),
```

```
    # but keep line 0 as the header.
```

```
    skip_list = range(1, current_skip) # can be large but 'engine="python"' can handle it
```

```
    df_chunk = pd.read_csv(
```

```
        CSV_FILE,
```

```
        skiprows=skip_list,
```

```
        nrows=chunk_size,
```

```
        usecols=wanted_columns,
```

```
        engine='python'
```

```
    )
```

```
# If there are no rows returned, we're done
```

```
if df_chunk.empty:
```

```
    break
```

```
# DATA CLEANING
```

```
# Clean all row contain NA value
```

```
df_cleaned = df_chunk.dropna(axis=0, how="any")
```

```
# Extract year
```

```
df_cleaned['calendar_last_scraped'] = pd.to_datetime(
```

```
    df_cleaned['calendar_last_scraped'],
```

```
    errors='coerce'
```

```
)
```

```
df_cleaned['year'] = df_cleaned['calendar_last_scraped'].dt.year
```

```

# Change the col from t/f to 1/0
df_cleaned['host_is_superhost'] = df_cleaned['host_is_superhost'].replace({'t': 1, 'f': 0})
df_cleaned['instant_bookable'] = df_cleaned['instant_bookable'].replace({'t': 1, 'f': 0})

all_chunks.append(df_cleaned)

# Advance the skip pointer by the chunk size
current_skip += chunk_size

<ipython-input-3-88ac68c5a396>:53: FutureWarning: Downcasting behavior in `replace` is deprecated and will be removed in :
df_cleaned['host_is_superhost'] = df_cleaned['host_is_superhost'].replace({'t': 1, 'f': 0})
<ipython-input-3-88ac68c5a396>:54: FutureWarning: Downcasting behavior in `replace` is deprecated and will be removed in :
df_cleaned['instant_bookable'] = df_cleaned['instant_bookable'].replace({'t': 1, 'f': 0})
<ipython-input-3-88ac68c5a396>:53: FutureWarning: Downcasting behavior in `replace` is deprecated and will be removed in :
df_cleaned['host_is_superhost'] = df_cleaned['host_is_superhost'].replace({'t': 1, 'f': 0})
<ipython-input-3-88ac68c5a396>:54: FutureWarning: Downcasting behavior in `replace` is deprecated and will be removed in :
df_cleaned['instant_bookable'] = df_cleaned['instant_bookable'].replace({'t': 1, 'f': 0})

df_merged = pd.concat([all_chunks[0], all_chunks[1]], axis=0, ignore_index=True)

# Furthur Data Cleaning
# Change host response rate
df_merged["host_response_rate"] = (
    df_merged["host_response_rate"]
    .astype(str) # convert to string
    .str.replace("%", "", regex=False) # remove the % sign
    .astype(float)
    .div(100) # convert 67 → 0.67
)

# Change cleaning fee
df_merged["cleaning_fee"] = (
    df_merged["cleaning_fee"]
    .fillna("") # fill missing with empty string
    .astype(str) # make sure everything is a string
    .str.replace("$", "", regex=False) # remove '$'
    .str.replace(",", "", regex=False) # remove commas
    .astype(float) # finally convert to float
)

# Compute compsite score
df_merged["composite_rating"] = df_merged[
    [
        "review_scores_accuracy",
        "review_scores_cleanliness",
        "review_scores_checkin",
        "review_scores_communication",
        "review_scores_location",
        "review_scores_value"
    ]
].mean(axis=1)

sia = SentimentIntensityAnalyzer()

# Fill and convert description to string
df_merged["description"] = df_merged["description"].fillna("").astype(str)

# Ensure the DataFrame has a simple integer index 0..N-1
df_merged = df_merged.reset_index(drop=True)


# Create the new column ahead of time (optional but recommended)
df_merged["sentiment_compound"] = None

# Loop over rows
for i in range(len(df_merged)):
    text = df_merged.loc[i, "description"]
    score = sia.polarity_scores(text)["compound"]
    df_merged.loc[i, "sentiment_compound"] = score

df_merged.to_csv('data.csv', index=True)

df = pd.read_csv('data.csv')
df.head()

```



Unnamed: 0		description	host_response_rate	host_is_superhost	cleaning_fee	calendar_last_scraped	review_scores_accurac
0	0	This home is perfect for families; aspiring ch...	0.67	0	100.0	2019-01-12	10
1	1	Our best memory foam pillows you'll ever sleep...	1.00	1	85.0	2019-01-11	10
2	2	This is a three story townhouse with the follo...	1.00	0	100.0	2019-01-12	9
3	3	A very Modern Hollywood Hills Zen style galler...	0.90	0	60.0	2019-01-12	8
4	4	Our distinctive bachelor's studio for 1-2 gues...	1.00	0	25.0	2019-01-11	8

EDA

```
# prompt: plot composite_rating vs host_response_rate

import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(10, 6))
sns.scatterplot(x='host_response_rate', y='composite_rating', data=df)
plt.title('Composite Rating vs. Host Response Rate')
plt.xlabel('Host Response Rate')
plt.ylabel('Composite Rating')
plt.show()

plt.figure(figsize=(10, 6))
sns.scatterplot(x='cleaning_fee', y='composite_rating', data=df)
plt.title('Composite Rating vs. Cleaning Fee')
plt.xlabel('Cleaning Fee')
plt.ylabel('Composite Rating')
plt.show()


plt.figure(figsize=(10, 6))
sns.scatterplot(x='sentiment_compound', y='composite_rating', data=df)
plt.title('Composite Rating vs. Sentiment Compound')
plt.xlabel('Sentiment Compound')
plt.ylabel('Composite Rating')
plt.show()

plt.figure(figsize=(10, 6))
sns.boxplot(x='host_is_superhost', y='composite_rating', data=df)
plt.title('Composite Rating vs. Host is Superhost')
plt.xlabel('Host is Superhost')
plt.ylabel('Composite Rating')
plt.show()


plt.figure(figsize=(10, 6))
sns.boxplot(x='instant_bookable', y='composite_rating', data=df)
plt.title('Composite Rating vs. Instant Bookable')
plt.xlabel('Instant Bookable')
plt.ylabel('Composite Rating')
plt.show()
```




Composite Rating vs. Host Response Rate

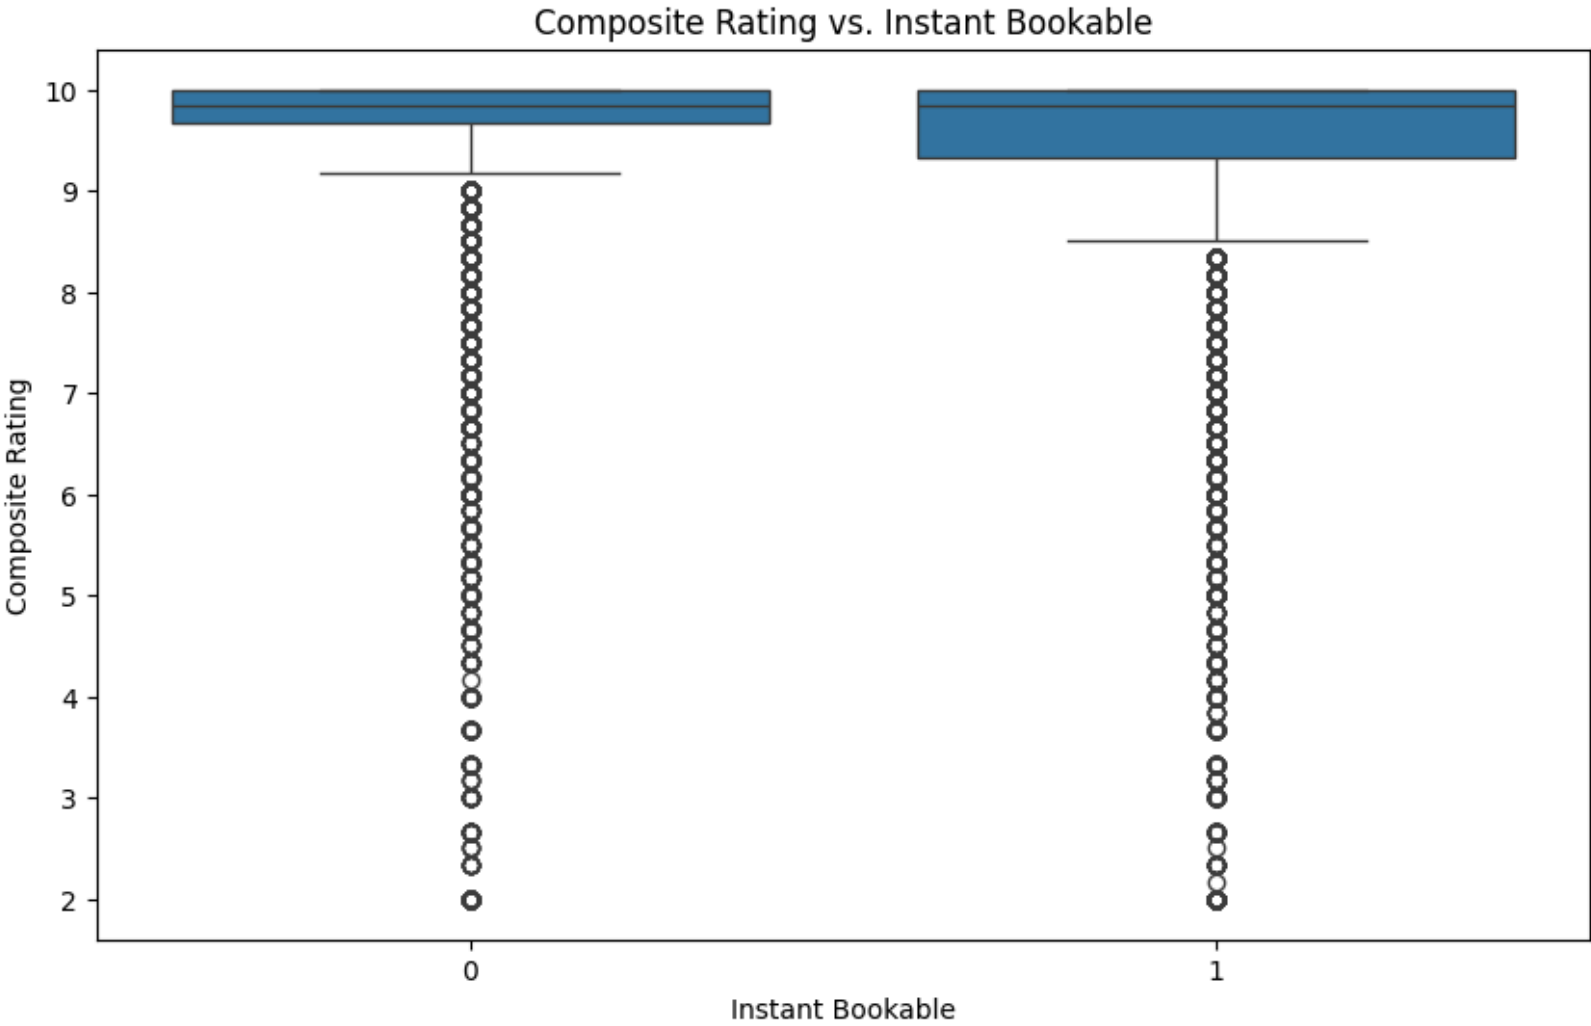
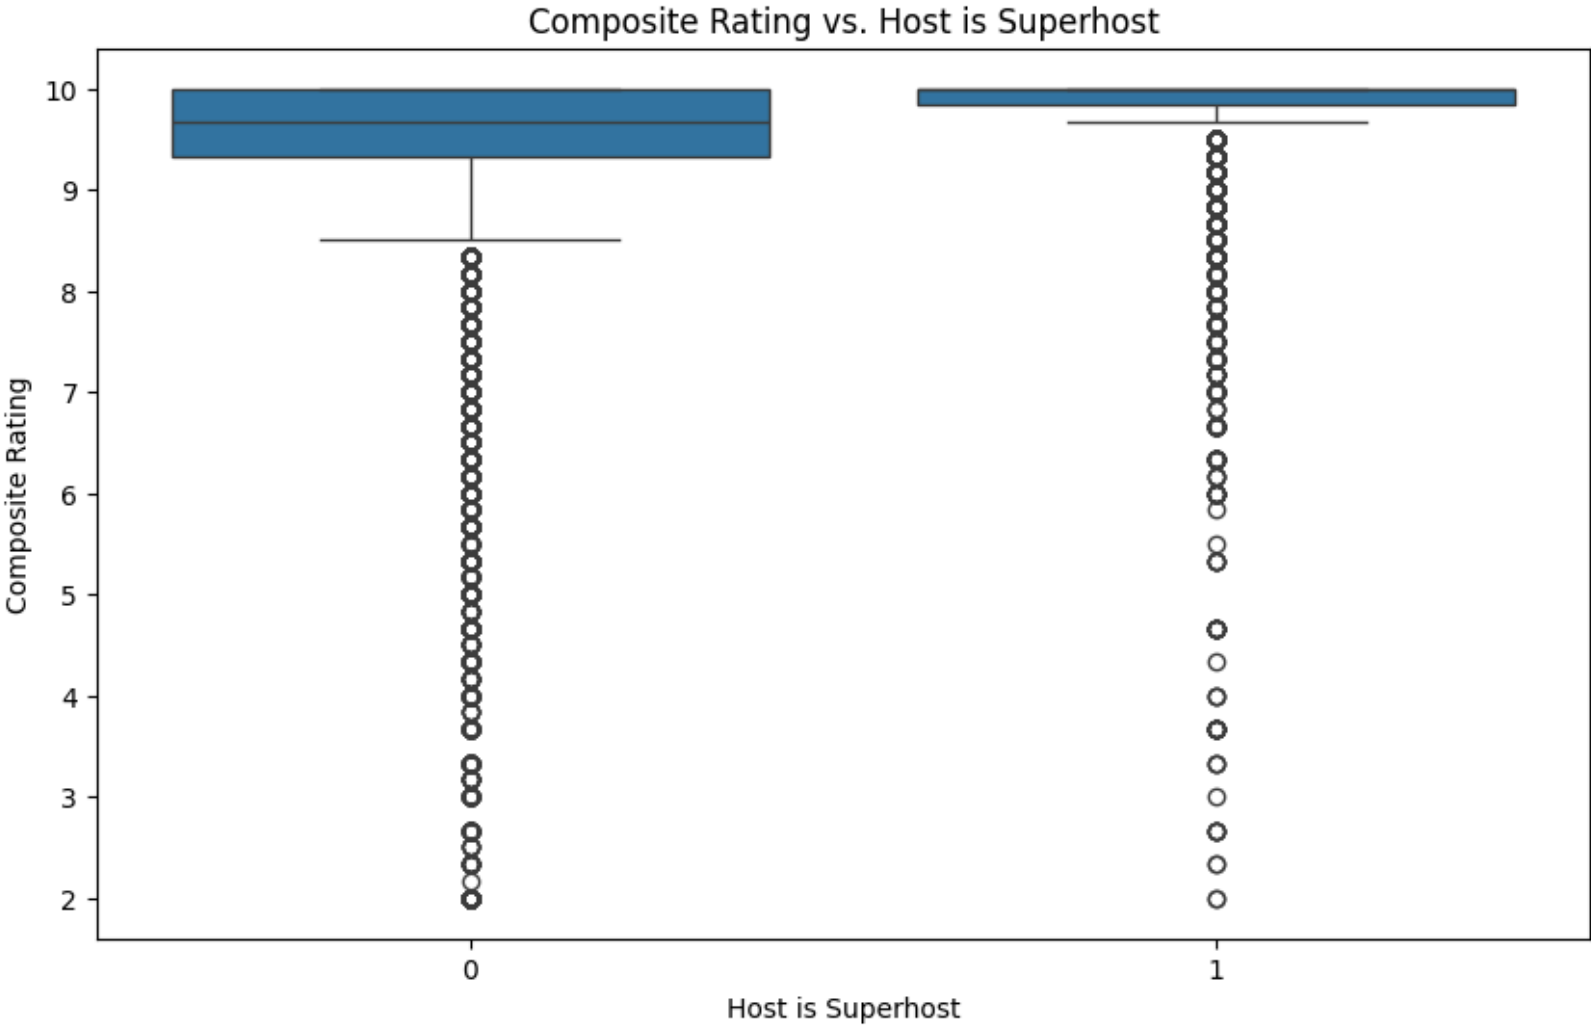



Composite Rating vs. Cleaning Fee



Composite Rating vs. Sentiment Compound







MLR Model

```
import statsmodels.api as sm

# Fit the multiple linear regression model
X = df[['host_response_rate', 'cleaning_fee', 'sentiment_compound', 'host_is_superhost', 'instant_bookable']]
X = sm.add_constant(X) # Add intercept
y = df['composite_rating']

mlr_model = sm.OLS(y, X).fit()

# Print the full summary (which includes the F-test result)
print(mlr_model.summary())

# Extract F-statistic and p-value
f_statistic = mlr_model.fvalue
p_value = mlr_model.f_pvalue

print(f"F-statistic: {f_statistic}")
print(f"P-value: {p_value}")

# Interpretation
if p_value < 0.05:
    print("The model is statistically significant (reject H0). At least one predictor explains variance in the dependent vari
else:
    print("The model is NOT statistically significant (fail to reject H0). The predictors collectively do not explain the out
```

→

OLS Regression Results						
=====						
Dep. Variable:	composite_rating	R-squared:	0.120			
Model:	OLS	Adj. R-squared:	0.120			
Method:	Least Squares	F-statistic:	1.302e+04			
Date:	Sun, 02 Mar 2025	Prob (F-statistic):	0.00			
Time:	06:56:58	Log-Likelihood:	-4.0869e+05			
No. Observations:	475581	AIC:	8.174e+05			
Df Residuals:	475575	BIC:	8.175e+05			
Df Model:	5					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]
-----						
const	9.1823	0.007	1336.962	0.000	9.169	9.196
host_response_rate	0.3013	0.007	45.967	0.000	0.288	0.314
cleaning_fee	-8.231e-05	9.88e-06	-8.335	0.000	-0.000	-6.3e-05
sentiment_compound	0.0838	0.003	25.622	0.000	0.077	0.090
host_is_superhost	0.3762	0.002	219.172	0.000	0.373	0.380
instant_bookable	-0.1240	0.002	-73.868	0.000	-0.127	-0.121
=====						
Omnibus:	479999.529	Durbin-Watson:	1.939			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	39180428.632			
Skew:	-4.917	Prob(JB):	0.00			
Kurtosis:	46.365	Cond. No.	1.35e+03			
=====						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.35e+03. This might indicate that there are strong multicollinearity or other numerical problems.

F-statistic: 13020.657353726368

P-value: 0.0

The model is statistically significant (reject H0). At least one predictor explains variance in the dependent variable.

```
from statsmodels.stats.outliers_influence import variance_inflation_factor
import pandas as pd
```